

Hierarchical Multi-Agent Repository Auditor

Interim Technical Report

By Amare Kassa February 25, 2026

1. Executive Summary

The Automaton Auditor is a production-grade, hierarchical multi-agent system that automatically audits software repositories and produces structured, evidence-backed technical assessments.

Rather than relying on a single large language model, the system decomposes auditing into specialized roles organized as a deterministic state graph:

- **Detectives** → collect factual evidence only
- **Judges** → independently reason over evidence
- **Chief Justice** → synthesize a final verdict

This separation enforces:

- factual correctness
- parallel execution
- deterministic state transitions
- reduced hallucination risk
- scalable architecture

The current implementation includes:

- typed shared state using Pydantic
 - AST-based static analysis
 - sandboxed repo execution
 - parallel detective layer
 - planned parallel judicial layer
 - structured JSON/Markdown audit reports
-

2. Architectural Philosophy

2.1 Core Principle: Separate Observation from Reasoning

A single agent performing both evidence gathering and judgment introduces:

- confirmation bias
- hallucinated facts
- non-reproducible behavior
- poor parallelism

We enforce a strict boundary:

Layer	Responsibility	Restrictions
Detective	Gather facts only	No scoring/opinion
Judicial	Score + argue	No filesystem/tools
Chief Justice	Aggregate + synthesize	No new facts

This mirrors courtroom-style adversarial reasoning and improves reliability.

3. System Overview (End-to-End Flow)

The full execution pipeline is shown below:

Data Flow & Failure Semantics

- Each stage consumes and emits typed state objects.
- Detective failures degrade gracefully (warnings) while judicial failures trigger retries or default scores.
- Critical failures (clone or workspace corruption) abort execution early to avoid misleading reports.

This design ensures:

- deterministic execution
- partial robustness
- safe degradation
- audit integrity

Two synchronization barriers exist:

1. Evidence aggregation
2. Judicial synthesis

This guarantees:

- reasoning never begins on partial evidence
- verdicts use complete information
- execution remains deterministic

Fig 1. End-to-End Flow

4. Key Architectural Decisions & Trade-offs

4.1 Typed State with Pydantic

Decision: Use strongly typed Pydantic models instead of raw dictionaries. **Why:**

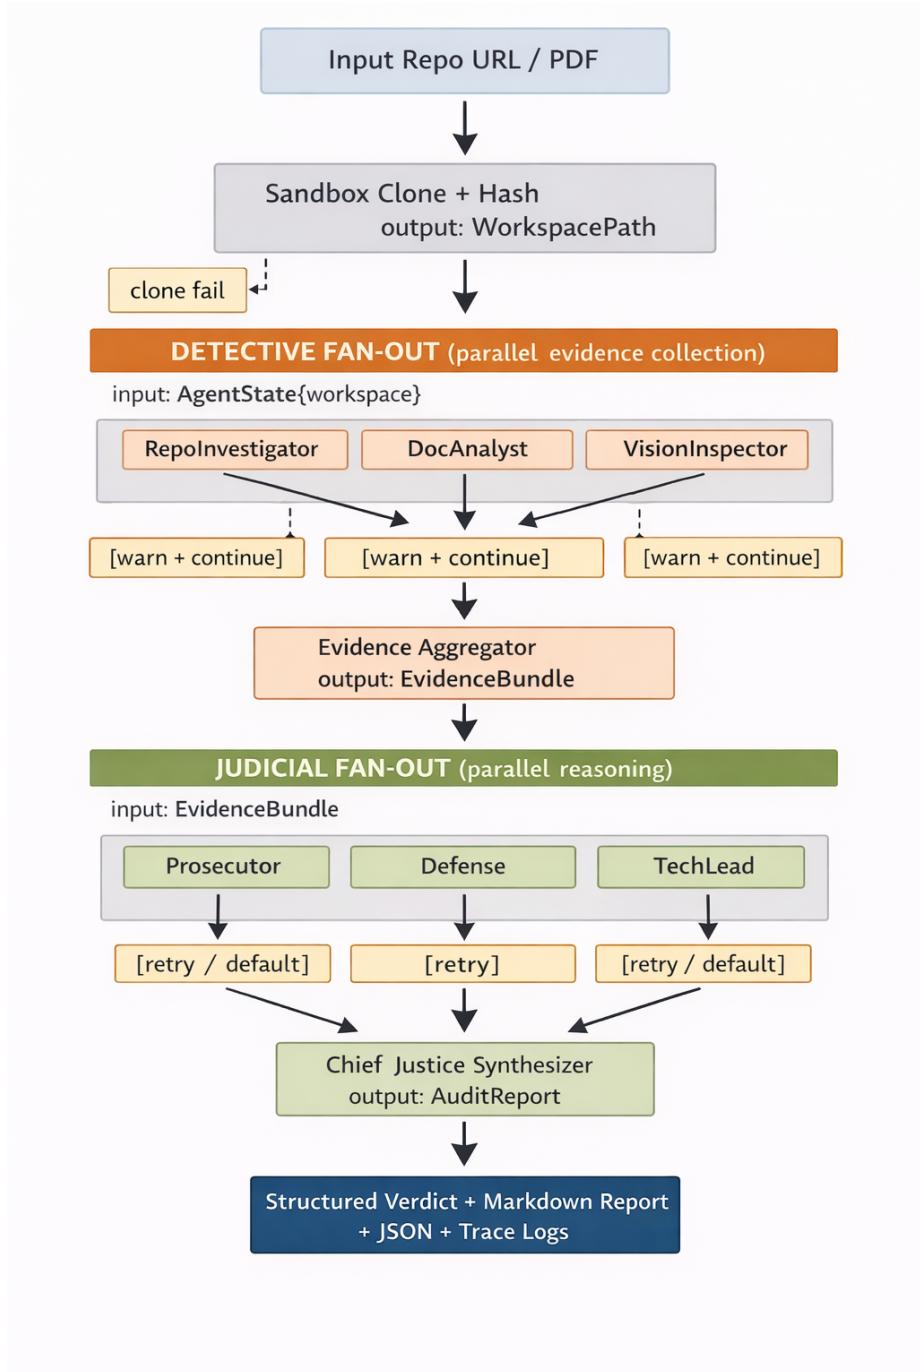


Figure 1: Automaton Auditor Flowchart

- validation guarantees
- schema enforcement
- safer multi-agent contracts
- easier serialization

Trade-offs:

Benefit	Cost
correctness	slightly slower
maintainability	more boilerplate
tooling support	stricter refactors

Rationale: Auditing is correctness-critical; safety outweighs micro-performance.

4.2 AST-Based Static Analysis

Decision: Use Python AST parsing instead of regex/grep. **Why:**

- syntactically correct extraction
- structured imports/exports
- language-aware reasoning

Trade-offs:

Benefit	Cost
accuracy	Python-specific
structure	higher complexity

Rationale: False positives from regex would undermine audit credibility.

4.3 Sandboxed Tool Execution

Decision: All repo operations run in isolated temporary directories. **Why:**

- prevents malicious repo execution
- avoids host contamination
- reproducible environments

Trade-offs:

Benefit	Cost
safety	small I/O overhead

Rationale: Security is non-negotiable when cloning arbitrary repos.

4.4 Parallel Fan-Out Architecture

Decision: Detectives and judges execute concurrently. **Why:**

- reduces latency
- increases independence
- enables horizontal scaling

Trade-offs:

Benefit	Cost
speed	coordination complexity

Rationale: Audits are embarrassingly parallel tasks.

5. Detective Layer Design

Detectives only gather evidence and do not reason.

Implemented agents:

- **RepoInvestigator** → Git metadata + file structure
- **DocAnalyst** → PDF parsing
- **VisionInspector** → diagram/image extraction

Properties:

- pure functions
- deterministic
- parallel safe
- side-effect isolated

Output: Dict [str, Evidence]

Each item includes:

- source
 - content
 - timestamp
 - metadata
-

6. Judicial Layer Design

Judges independently score evidence using rubric criteria.

Planned agents:

- **Prosecutor** → find risks

- **Defense** → find strengths
- **TechLead** → assess engineering quality

Why independent judges?

Prevents:

- single-model bias
- cascading hallucination
- overconfident conclusions

Produces:

- dialectical reasoning
- explainable results

Output: JudicialOpinion[]

Each opinion includes:

- criterion
 - score
 - argument
 - cited evidence IDs
-

7. Synthesis Engine (Chief Justice)

The Chief Justice:

- aggregates scores
- averages per criterion
- composes summary
- generates report

Produces:

- JSON verdict
- Markdown audit report

This stage is deterministic and contains no new reasoning.

8. Known Gaps & Forward Plan

Remaining Work:

- **Judicial implementation:** prompt templates, rubric scoring, parallel execution
- **Aggregation:** score normalization, conflict handling, weighted scoring
- **Reporting:** structured markdown, evidence citations, remediation plans
- **Observability:** trace logging, timing metrics, failure tracking

9. Anticipated Failure Modes & Mitigations

Failure	Risk	Mitigation
malicious repos	code execution	sandbox clone
hallucinated evidence	false claims	detective-only facts
judge bias	skewed scoring	multiple judges
race conditions	incomplete data	fan-in barrier
large repos	timeouts	batching
PDF parsing errors	missing docs	fallback logs

10. Scalability Strategy

The architecture scales horizontally:

- add more detectives for new modalities
- add more judges for deeper reasoning
- execution time \approx slowest node only

No redesign required.

11. Conclusion

The Automaton Auditor implements a structured, reliable approach to repository auditing using:

- typed shared state
- deterministic graph orchestration
- parallel evidence gathering
- adversarial reasoning
- explainable verdict synthesis

By explicitly separating observation from judgment and enforcing synchronization barriers, the system prioritizes correctness, reproducibility, and scalability over convenience.

Next Milestone: complete the judicial layer and reporting engine to deliver fully automated, production-grade audits.