

# Interim Technical Report – Automaton Auditor

By Amare Kassa February 25, 2026

## 1. Introduction

- Objective: Build a hierarchical multi-agent auditor (Detectives → Judges → ChiefJustice) for Week 2.
  - Scope of interim submission: foundation of Detective Layer, state modeling, repo/PDF/AST tools, and architecture planning for Judicial Layer.
- 

## 2. Architecture Decisions So Far

### Typed Agent State

- **Pydantic vs. Python dicts:**
  - Enforces strict type validation.
  - Ensures structured AgentState across nodes and fan-out/fan-in.
  - Prevents accidental mutations and supports traceable JSON serialization.

### StateGraph Orchestration

- Nodes (`Node` class) with `run()` function for typed state transition.
- Supports sequential and conditional execution.
- Fan-out/fan-in planned for Detectives → EvidenceAggregator.

### Detective Layer Tools

- **AST Parsing (`ast_tools`):** Extract exports, imports, and function signatures from Python source.
- **Git Tools (`git_tools`):** Clone, list commits, enumerate tracked files.
- **PDF Tools (`pdf_tools`):** Extract, chunk text for DocAnalyst querying.
- **Filesystem Utilities (`filesystem.py`):** Temp directories, safe file writes, SHA256 hashes.

### Sandboxing Strategy

- Clone and analyze repos in temporary directories using `tempfile`.
  - Prevents accidental overwrites, ensures isolated environment per repo.
- 

## 3. Known Gaps

- **Judicial Layer:**
  - Personas (Prosecutor, Defense, TechLead) not yet fully implemented.

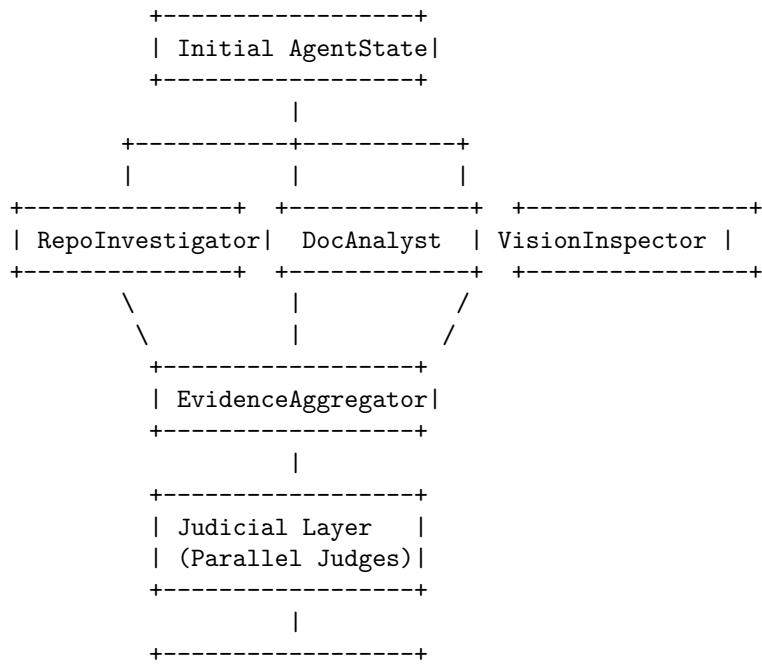
- Structured JSON outputs with criterion-specific reasoning pending.
  - **ChiefJustice / Synthesis Engine:**
    - Conflict resolution and score synthesis logic not yet implemented.
    - Audit report generation in Markdown still a placeholder.
  - **Vision Inspector:**
    - Multimodal diagram analysis optional, not yet integrated.
- 

#### 4. Planned Implementation

1. **Judicial Layer**
    - Load evidence from EvidenceAggregator.
    - Execute each Judge in parallel using distinct prompts/system logic.
    - Generate `JudicialOpinion` objects for each rubric criterion.
  2. **ChiefJusticeNode / Synthesis**
    - Aggregate opinions.
    - Apply rules: `security_override`, `fact_supremacy`, `dissent_requirement`.
    - Generate structured Markdown Audit Report with Executive Summary and Remediation Plan.
- 

#### 5. Planned Diagrams

##### StateGraph Flow (Detective Fan-Out / Fan-In)



```
| ChiefJusticeNode |
+-----+
|           |
+-----+
| Audit Report PDF |
+-----+
```

---

## 6. Next Steps

- Implement parallel Judge execution with persona-specific structured output.
- Integrate ChiefJustice conflict resolution and synthesis.
- Generate Markdown/structured audit reports.
- Optional: VisionInspector multimodal analysis.
- Batch-processing multiple repositories for adversarial MinMax testing.